model completes during the reverse process. This allows us to interpolate at both fine granularities and coarse granularities. In the limiting case of 0 diffusion steps, the interpolation mixes source images in pixel space. On the other hand, after 1000 diffusion steps, source information is lost and interpolations are novel samples.

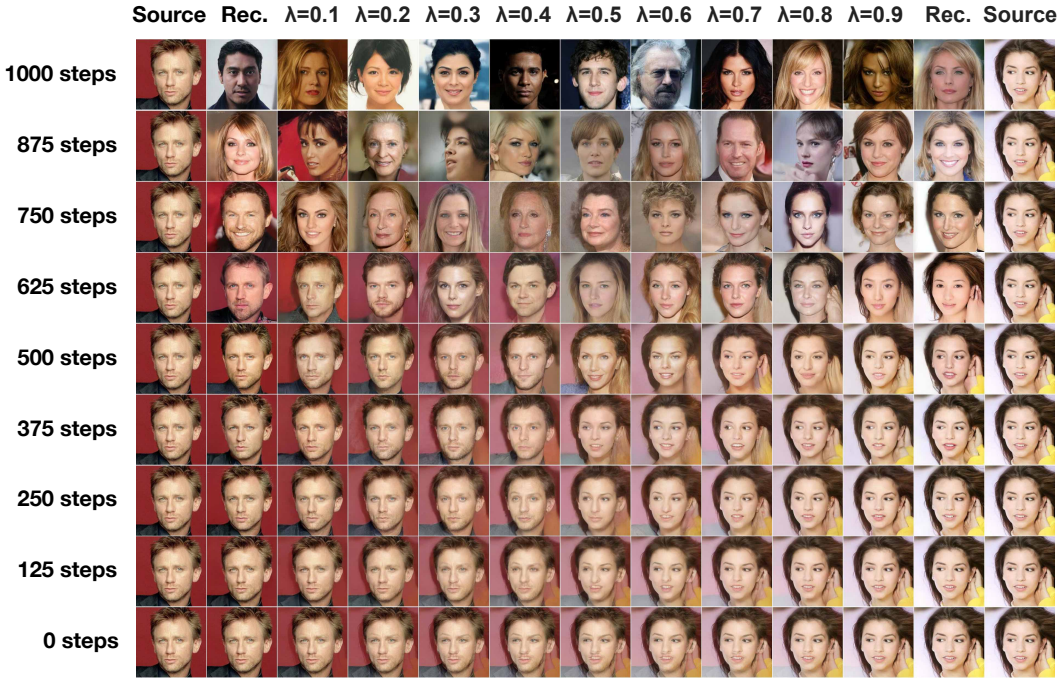

Figure 9: Coarse-to-fine interpolations that vary the number of diffusion steps prior to latent mixing.

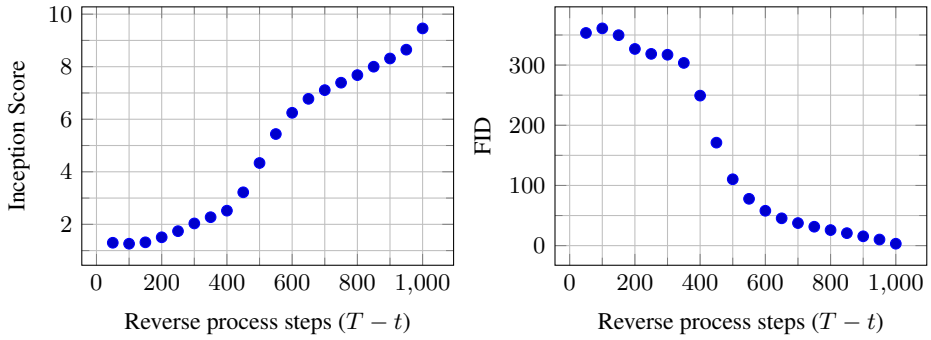

Figure 10: Unconditional CIFAR10 progressive sampling quality over time

Figure 11: CelebA-HQ 256×256 generated samples

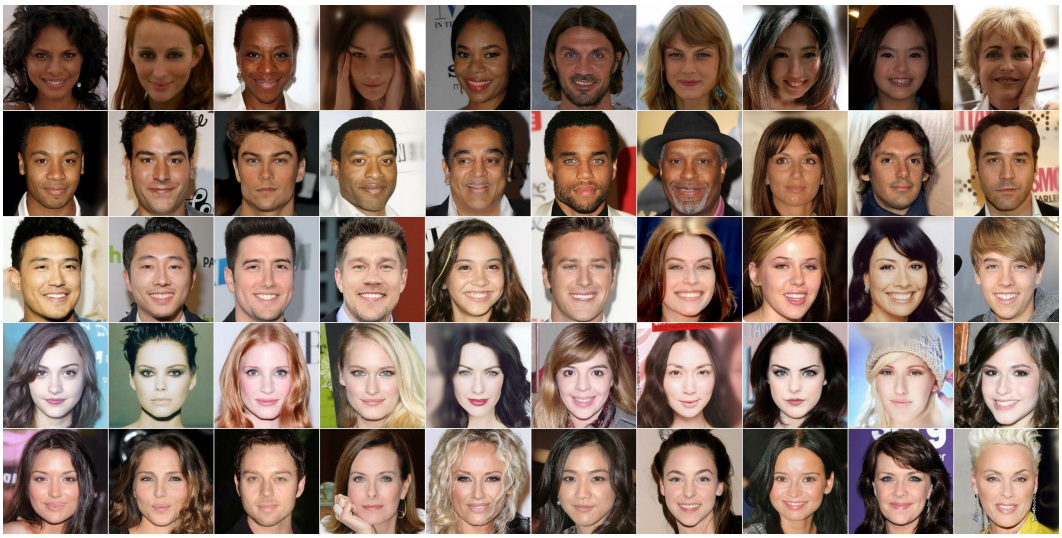

(a) Pixel space nearest neighbors

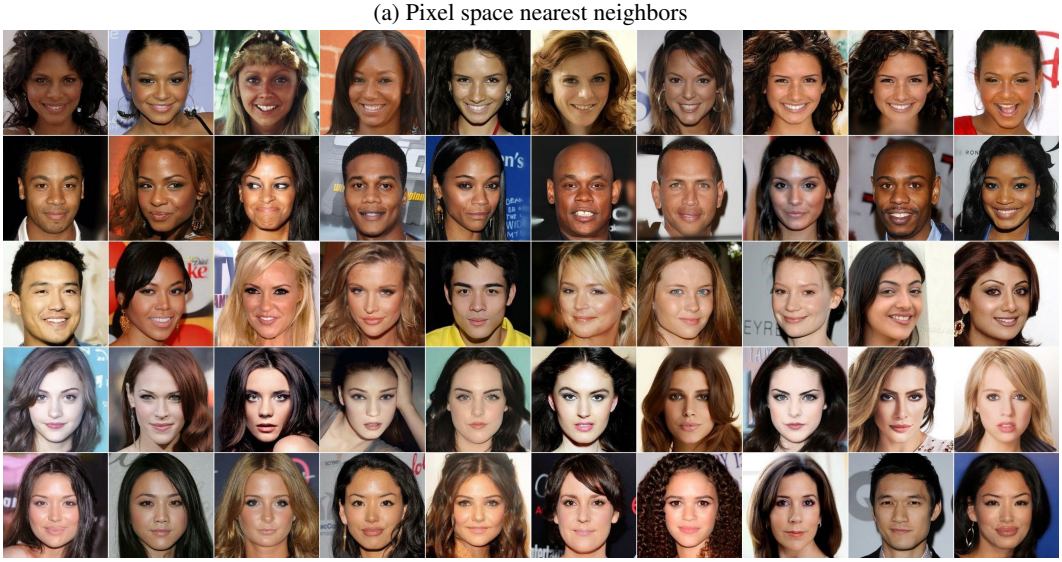

(b) Inception feature space nearest neighbors

Figure 12: CelebA-HQ 256×256 nearest neighbors, computed on a 100×100 crop surrounding the faces. Generated samples are in the leftmost column, and training set nearest neighbors are in the remaining columns.